

MATH 250: Mathematical Data Visualization

Singular value decomposition and principal components analysis

Peter A. Gao

2025-02-26

San José State University

Warm-up: Ranks of common flags

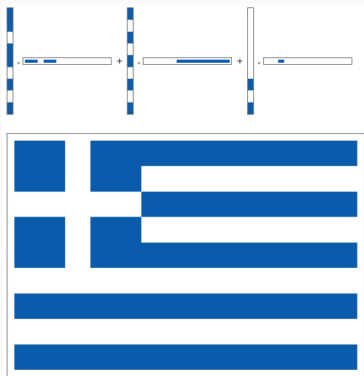


Figure 1: Rank-1 decomposition of Greek Flag by Michael E2

What would the rank of Denmark's flag be?



Singular Value Decomposition

Any matrix $A \in \mathbb{R}^{m \times n}$ can be factorized

$$A = U\Sigma V^T$$

with

- $U \in \mathbb{R}^{m \times m}$ an orthogonal matrix of the **left singular vectors** of A
- $\Sigma \in \mathbb{R}^{m \times n}$ an $m \times n$ **singular value** matrix
- $V \in \mathbb{R}^{n \times n}$ an orthogonal matrix of the **right singular vectors** of A

Review: SVD

The columns of U form an orthonormal basis for the **column space** of A while the columns of V (rows of V^\top) form an orthonormal basis for the **row space**.

The key property to remember for the singular vectors:

$$A\mathbf{v}_i = \sigma_i\mathbf{u}_i, \quad i = 1, \dots, r$$

Where \mathbf{v}_i and \mathbf{u}_i are the i th right and left singular vectors, respectively, σ_i is the i th singular value, and r is the rank of A .

Example:

$$\begin{bmatrix} 3 & 0 \\ 4 & 5 \end{bmatrix} = \begin{bmatrix} 1/\sqrt{10} & -3/\sqrt{10} \\ 3/\sqrt{10} & 1/\sqrt{10} \end{bmatrix} \begin{bmatrix} 3\sqrt{5} & 0 \\ 0 & \sqrt{5} \end{bmatrix} \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ -1/\sqrt{2} & 1/\sqrt{2} \end{bmatrix}$$

Decomposing linear transformations

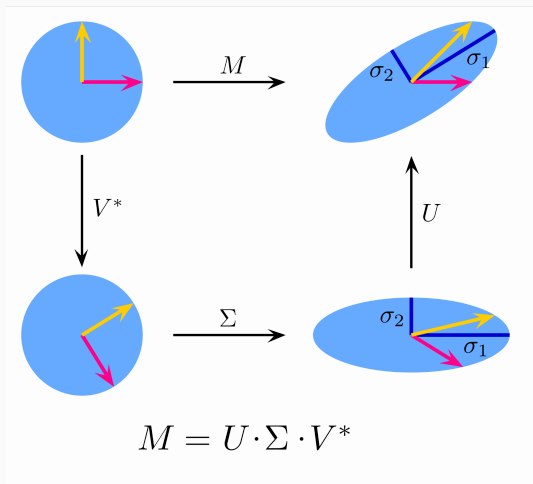


Figure 2: Geometric interpretation of SVD, by Georg-Johann

Review: SVD

In other words:

$$AV = U\Sigma \iff A \begin{bmatrix} \mathbf{v}_1 & \cdots & \mathbf{v}_n \end{bmatrix} = \begin{bmatrix} \mathbf{u}_1 & \cdots & \mathbf{u}_m \end{bmatrix} \left[\begin{array}{ccc|c} \sigma_1 & & & 0 \\ & \ddots & & \\ & & \sigma_r & 0 \\ \hline & 0 & & 0 \end{array} \right]$$

We can also write A in a reduced SVD form:

$$AV_r = U_r \Sigma_r$$

making Σ_r a diagonal $r \times r$ matrix and removing the last singular vectors from V and U .

Review: SVD as sum of rank-1 matrices

We can also write A as a sum of rank-1 matrices:

$$A = \sigma_1 \mathbf{u}_1 \mathbf{v}_1^\top + \cdots + \sigma_r \mathbf{u}_r \mathbf{v}_r^\top$$

and obtain the best rank- k approximation:

$$A_k = \sigma_1 \mathbf{u}_1 \mathbf{v}_1^\top + \cdots + \sigma_k \mathbf{u}_k \mathbf{v}_k^\top$$

Theorem

(Eckart-Young) If B has rank k , then $\|A - A_k\| \leq \|A - B\|$ in either the Frobenius or ℓ^2 norm.

Review: Other SVD properties

- V contains orthonormal eigenvectors of $A^T A$ and U contains orthonormal eigenvectors of AA^T .
- $\sigma_1^2, \dots, \sigma_r^2$ are the nonzero eigenvalues of both $A^T A$ and AA^T .
- If S is symmetric positive definite, $U\Sigma V^T = Q\Lambda Q^T$.
- \mathbf{v}_1 maximizes $\|A\mathbf{x}\|/\|\mathbf{x}\|$, achieving a value of σ_1 .

Application: Image compression

Original (2419 kb)



Figure 3: Stephan's quintet

Application: Image compression

Each pixel is a value from 0-255 representing a color from white to black.

We can thus treat this image as a matrix, compute its SVD and the best rank- k approximation.

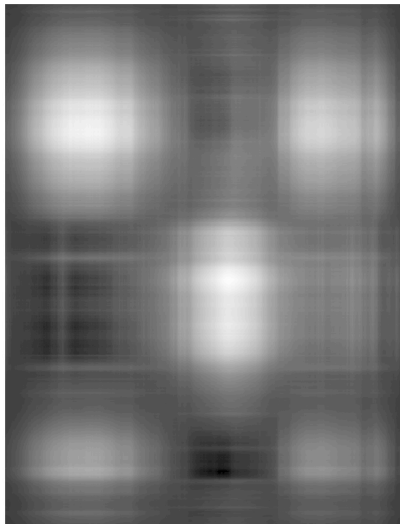
Plotting the rank- k approximation yields a compressed version of our original image.

Application: Image compression}

Original (2419 kb)



Rank-2 approximation (11 kb)

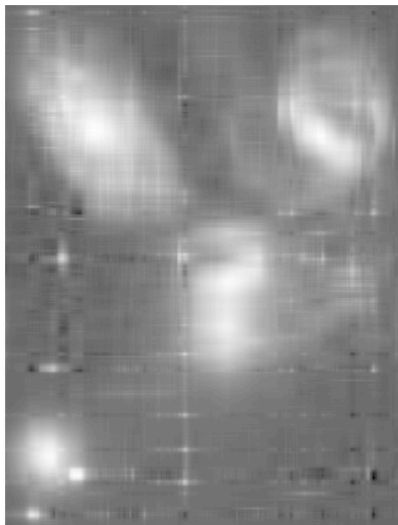


Application: Image compression

Original (2419 kb)



Rank-10 approximation (55.2 kb)

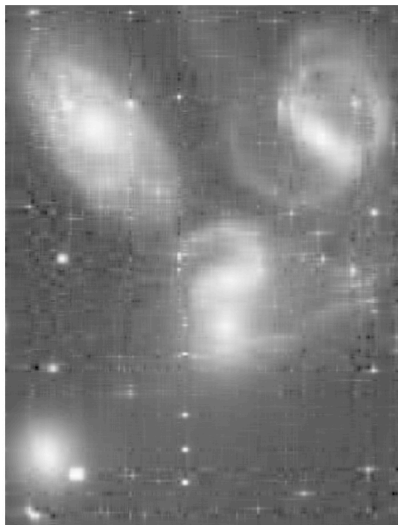


Application: Image compression

Original (2419 kb)



Rank-20 approximation (110.4 kb)



Application: Image compression

Original (2419 kb)



Rank-50 approximation (276 kb)



Application: Image compression

The SVD provides a crude approach to image compression, which is the “best” in the sense that it minimizes the matrix distance between these images.

However, when viewing two images, this may not be the right “distance” to be using.

When noise has been added to our image, the SVD can also be used to denoise and clean up images.

USPS handwritten digits data:

- 9298 16 x 16 images of handwritten digits, split into training and test datasets.
- Centered and scaled to be the same size.

Example: handwritten digits

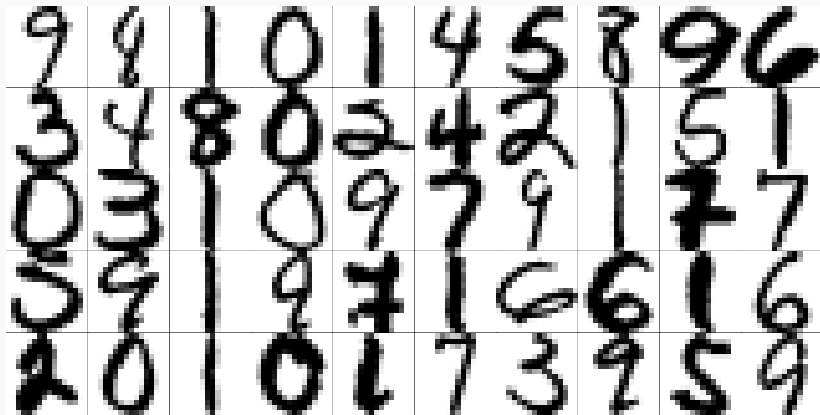


Figure 4: Handwritten 16 x 16 digits from USPS dataset Hull (1994)

Example: handwritten digits

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]	[,15]	[,16]
[1,]	0	0	0	0	0	11	167	197	29	0	0	0	0	0	0	0
[2,]	0	0	0	0	22	207	255	204	0	0	0	0	0	0	0	0
[3,]	0	0	0	95	248	255	160	7	0	0	0	0	0	0	0	0
[4,]	0	58	175	255	255	226	117	146	117	88	88	29	0	0	0	0
[5,]	84	255	255	255	204	145	145	204	145	174	229	255	197	80	0	0
[6,]	32	65	36	7	0	0	0	0	0	0	3	160	255	255	65	0
[7,]	0	0	0	0	0	0	0	0	0	0	0	204	255	236	21	0
[8,]	0	0	0	0	0	0	0	0	0	59	175	255	225	40	0	0
[9,]	0	0	0	0	0	0	22	110	226	255	233	145	0	0	0	0
[10,]	0	0	22	132	190	219	255	255	255	255	190	132	73	0	0	0
[11,]	0	0	7	101	130	72	14	14	14	72	101	159	251	212	37	0
[12,]	0	0	0	0	0	0	0	0	0	0	0	0	25	255	255	44
[13,]	0	0	0	0	0	0	0	0	0	0	0	0	26	255	255	101
[14,]	0	0	116	95	0	0	0	0	0	0	0	44	193	255	247	0
[15,]	0	0	0	138	154	37	37	37	66	125	212	255	236	130	21	0
[16,]	0	0	0	0	50	108	166	196	196	196	137	79	10	0	0	0

Figure 5: Handwritten digit from USPS dataset, as matrix

The “typical” digits

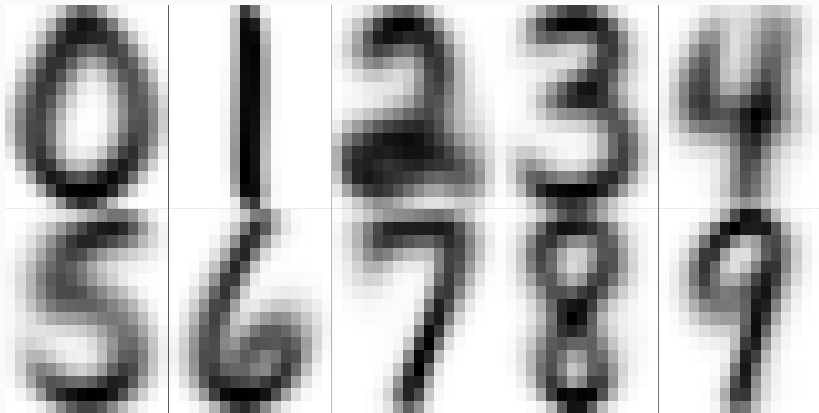


Figure 6: Centroids from USPS dataset (element-wise means)

A naive classification algorithm

1. For each new image i , calculate its distance from the centroids 0-9.
2. Label the new image i based on the closest centroid.

This achieves 75% accuracy. Note the work needed to create all these black-and-white, centered images.

Creating SVD-based representations of each digit

Let n_i be the number of images of digit i in the training set. For each digits, construct a $n_i \times 256$ matrix:

$$\begin{array}{ccc} n_i \text{ rows} & \boxed{A} & \\ & & 256 \text{ columns} \end{array}$$

The right singular vectors \mathbf{v}_i of A form an orthonormal basis in the space of images.

For a given digit, the first few singular vectors can be used to reconstruct each image in the training set.

Creating SVD-based representations of each digit

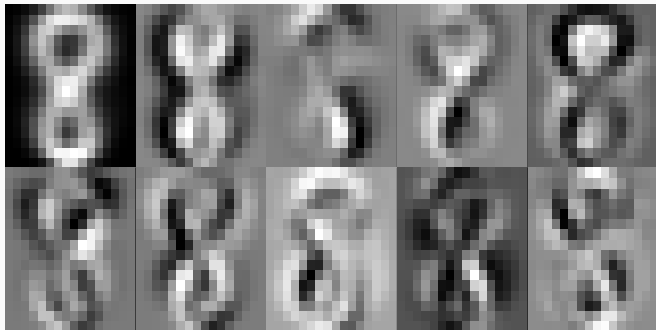


Figure 7: First 10 basis images for USPS 8 digits

Creating SVD-based representations of each digit

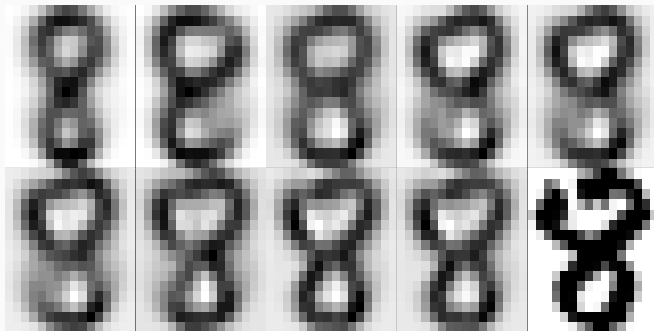


Figure 8: A single 8 image (lower right), projected on the first k basis images, for $k = 1$ to $k = 9$

Creating SVD-based representations of each digit



Figure 9: A single 8 image (lower right), projected on the first k basis images for the digit 9, for $k = 1$ to $k = 9$

SVD basis classification

1. For each new image i , calculate its representation in the SVD basis for each digit.
2. Label the new image i based on the most accurate representation.

For an unknown image \mathbf{z} , we can approximate it in a basis using a least squares solution:

$$\min_{\mathbf{c}} \left\| \mathbf{z} - \sum_{i=1}^k c_i \mathbf{u}_i \right\|$$

We can repeat this process for each basis for 0-9 and identify the digit that yields the most accurate representation.

The accuracy of this classification method depends on the dimension k of the basis:

# basis images	1	2	4	6	8	10
accuracy	80	86	90	90.5	92	93

Figure 10: Accuracy by number of basis images (Elden 2007)

Principal components analysis

You may often hear “PCA is just SVD.” It is—sort of.

SVD

- a matrix method
- $m \times n$ matrix

PCA

- a data analysis method
- $n \times p$ data matrix

Let's start with PCA and show how it relates to the SVD.

Definition

The **covariance** between two random variables X and Y is defined as

$$\text{Cov}(X, Y) = E[(X - E(X))(Y - E(Y))]$$

Remark

If \mathbf{x} is a p -dimensional random vector, its **covariance matrix** is defined to be

$$\text{Cov}(\mathbf{x}) = E[(\mathbf{x} - E(\mathbf{x}))(\mathbf{x} - E(\mathbf{x}))^\top]$$

Thus, the covariance matrix is positive semidefinite.

$$\begin{aligned}\text{Cov}(\mathbf{x}) &= E[(\mathbf{x} - E(\mathbf{x}))(\mathbf{x} - E(\mathbf{x}))^\top] \\ &= \begin{bmatrix} \text{Var}(X_1) & \text{Cov}(X_1, X_2) & \cdots & \text{Cov}(X_1, X_p) \\ \text{Cov}(X_1, X_2) & \text{Var}(X_2) & \cdots & \text{Cov}(X_2, X_p) \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}(X_1, X_p) & \text{Cov}(X_2, X_p) & \cdots & \text{Var}(X_p) \end{bmatrix}\end{aligned}$$

Remark

As a result,

$$\text{Cov}(A\mathbf{x} + \mathbf{b}) = A[\text{Cov}(\mathbf{x})]A^\top$$

Reviewing statistics

If we have a sample of iid random vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \sim \mathbf{x}$, we can combine them into an $n \times p$ data matrix:

$$X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

Definition

The **sample covariance** of \mathbf{x} is defined as

$$\widehat{S} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^\top = \frac{1}{n} X_c^\top X_c$$

where $\bar{\mathbf{x}}$ is the sample mean and X_c is a centered version of X (Exercise).

Motivation

We want to project our p -dimensional data into a simpler q -dimensional space. We will try to choose the “most important” q dimensions along which the data have maximum variance.

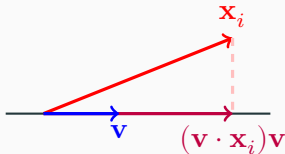
We can start with the example where $q = 1$. What does it mean to find the “optimal” one-dimensional projection of our data? Assume all of our data is centered, so the mean of each column of X is zero.

PCA: One-dimensional case

Idea (from Shalizi)

Choose the unit vector \mathbf{v} such that when we project our data vectors $\mathbf{x}_1, \dots, \mathbf{x}_n$ onto \mathbf{v} , the residual error is minimized:

$$\text{MSE}(\mathbf{v}) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - (\mathbf{v} \cdot \mathbf{x}_i) \mathbf{v}\|^2$$



Idea (from Shalizi)

This turns out to be equivalent to maximizing the sample variance of lengths of the projections onto \mathbf{v} (since the columns of X are centered):

$$\begin{aligned}\widehat{\text{Var}}(\mathbf{v} \cdot \mathbf{x}_i) &= \frac{1}{n} \sum_{i=1}^n (\mathbf{v} \cdot \mathbf{x}_i)^2 \\ &= \frac{1}{n} (X\mathbf{v})^\top (X\mathbf{v}) \\ &= \frac{1}{n} \mathbf{v}^\top X^\top X \mathbf{v} \\ &= \mathbf{v}^\top \widehat{S} \mathbf{v}\end{aligned}$$

In other words, we are simply maximizing the Rayleigh quotient $\mathbf{v}^\top \hat{S} \mathbf{v}$. How do we find the maximizing \mathbf{v} ?

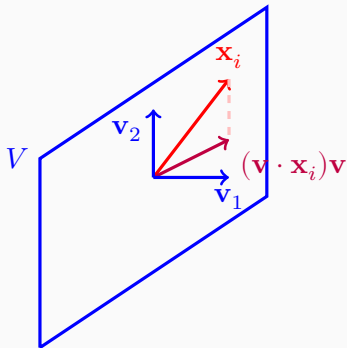
In other words, we are simply maximizing the Rayleigh quotient $\mathbf{v}^\top \hat{S} \mathbf{v}$. How do we find the maximizing \mathbf{v} ?

From last week, the maximizing \mathbf{v} is the eigenvector of \hat{S} with the largest eigenvalue λ_1 .

Thus, $\mathbf{v}^\top \hat{S} \mathbf{v}$ achieves maximum value λ_1 .

PCA: Multi-dimensional case

For $q > 1$, we can generalize our approach. Instead of the single vector along which the projected data has maximum variance, we are looking for a k -dimensional plane along which our projected data has maximum variance.



Theorem

The q -dimensional plane along which our projected data has maximum variance has an orthonormal basis given by the first q eigenvectors of \hat{S} and the total variance of the projections is given by $\lambda_1 + \dots + \lambda_k$.

Computing $X^T X$ is potentially expensive and can lead to an ill-conditioned matrix.

Luckily, the first q eigenvectors of $X^T X$ are also given by...

Computing $X^T X$ is potentially expensive and can lead to an ill-conditioned matrix.

Luckily, the first q eigenvectors of $X^T X$ are also given by... the first q right singular vectors of X (remember, X is centered). The variance captured by the q -dimensional projection plane is $\lambda_1 + \dots + \lambda_q = \sigma_1^2 + \dots + \sigma_q^2$.

Principal components analysis typically involves identifying a set of maximum-variance directions

$$V_q = [\mathbf{v}_1 \quad \cdots \quad \mathbf{v}_q] \in \mathbb{R}^{p \times q}$$

and the corresponding coordinates of each of the observations in the new basis

$$Y_q = [\mathbf{y}_1 \quad \cdots \quad \mathbf{y}_q] \in \mathbb{R}^{n \times q}$$

where

$$Y_q = XV_q = U_q \Sigma_q$$

Definitions

1. **Principal directions:** The unit eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_q$ that form a basis for the q -dimensional subspace on which the projected data have maximum variance.
2. **Variable loadings:** Usually also referring to the eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_q$ and in particular the scalar components. For example, v_{11} , the first component of \mathbf{v}_1 represents the contribution (loading) of \mathbf{x}_1 to the first principal direction.
3. **Principal components:** The columns of the rotated data matrix XV_q representing the new coordinates after projected the data onto the q -dimensional subspace spanned by the principal directions.

We can say

- The unit eigenvector \mathbf{v}_j is the j th **principal direction** of the data;
- The **principal components** $Y \in \mathbb{R}^{n \times q}$ are the coefficients obtained by projecting X on the first q **principal directions** of the data.

In essence, PCA is a change of coordinate system, where the new axes are the principal directions/axes of the data and the new coordinates the principal components. These principal directions are orthogonal.

How many principal components?

If $n \geq p$, then as long as the columns of X are linearly independent, the rank of X is p , so we can have up to p principal components/directions.

What if $n < p$? Still feasible and as long as the columns of X are linearly independent, but the maximum number of principal components is equal to n .

Geometric interpretation

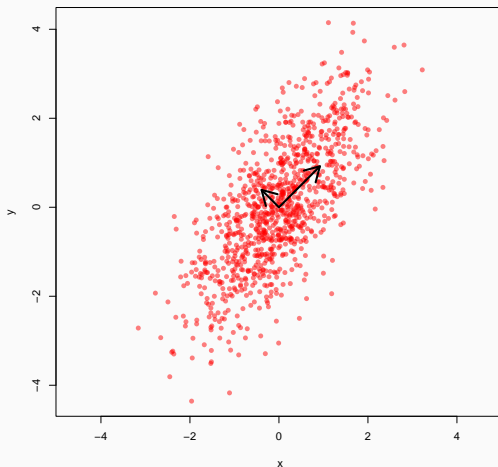


Figure 11: Example data with **principal directions**

Geometric interpretation

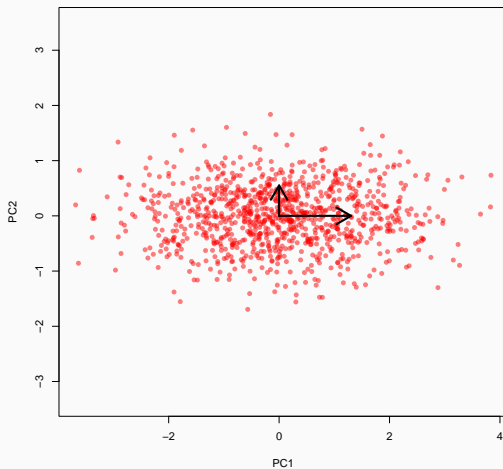


Figure 12: Example principal components

Given a data matrix $X \in \mathbb{R}^{n \times p}$ and an integer q ,

1. Center X by subtracting out the mean, if necessary.
2. Carry out rank- q SVD on $X \approx U_q \Sigma_q V_q^\top$ (using an efficient algorithm)
3. Compute the principal components $Y = U_q \Sigma_q$.



In honor of San Jose's new soccer team Bay FC, we'll take a look at data from the National Women's Soccer League, using the `nws1R` package.

In 2023, there were 12 teams. We can download team-level data for each team including variables like **goals**, **assists**, and **goals allowed**.

Application: NWSL Teams

```
library(nwslR)
library(tidyverse)
teams <- load_teams() |>
  filter(last_season >= 2023 & first_season <= 2023 & team_abbreviation != "UTA")
nwsl23 <-
  lapply(
    teams$team_abbreviation,
    function(x) load_team_season_stats(team_id = x, season = "2023")
  ) |>
  bind_rows()

nwsl23_mat <- nwsl23[, c(-1, -2)] |>
  dplyr::select("possession_pct",
               "goals", "assists", "pass_pct",
               "goal_conversion_pct", "clean_sheets",
               "shot_accuracy", "shots_total", "goals_conceded",
               "tackled") |>
  as.matrix(nrow = nrow(nwsl23))
rownames(nwsl23_mat) <-
  teams$team_abbreviation[match(nwsl23$team_id, teams$team_id)]
```

Application: NWSL Teams

```
head(nwsl23_mat)
```

	possession_pct	goals	assists	pass_pct	goal_conversion_pct	clean_sheets
CHI	47	19	11	75.15	17.59	3
HOU	48	10	5	71.08	7.87	6
NJY	53	17	8	73.67	11.26	4
RGN	49	23	18	73.97	17.04	4
ORL	47	15	9	75.07	10.20	4
POR	53	34	27	77.53	15.25	4

	shot_accuracy	shots_total	goals_conceded	tackled
CHI	46.30	108	33	300
HOU	44.88	127	12	259
NJY	48.34	151	14	303
RGN	47.41	135	18	220
ORL	47.62	147	21	314
POR	52.02	223	21	274

Application: NWSL Teams

```
pca <- prcomp(nwsl23_mat, scale = T)
options(width = 60)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5
Standard deviation	1.9471	1.5613	1.1729	1.1076	0.84948
Proportion of Variance	0.3791	0.2438	0.1376	0.1227	0.07216
Cumulative Proportion	0.3791	0.6229	0.7605	0.8832	0.95532

	PC6	PC7	PC8	PC9
Standard deviation	0.50289	0.36546	0.20599	0.13384
Proportion of Variance	0.02529	0.01336	0.00424	0.00179
Cumulative Proportion	0.98061	0.99396	0.99821	1.00000

	PC10
Standard deviation	0.002521
Proportion of Variance	0.000000
Cumulative Proportion	1.000000

Application: NWSL Teams

```
str(pca) # structure of pca object
```

List of 5

```
$ sdev      : num [1:10] 1.947 1.561 1.173 1.108 0.849 ...
$ rotation: num [1:10, 1:10] 0.368 0.422 0.415 0.386 0.337 ...
.- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:10] "possession_pct" "goals" "assists" "pass_pct" ...
.. ..$ : chr [1:10] "PC1" "PC2" "PC3" "PC4" ...
$ center   : Named num [1:10] 49.9 19.7 12.2 74.6 13.9 ...
.- attr(*, "names")= chr [1:10] "possession_pct" "goals" "assists" "pass_pct"
$ scale     : Named num [1:10] 3.63 5.71 5.91 3.03 3.19 ...
.- attr(*, "names")= chr [1:10] "possession_pct" "goals" "assists" "pass_pct"
$ x         : num [1:12, 1:10] -1.046 -2.551 -0.679 0.891 -1.577 ...
.- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:12] "CHI" "HOU" "NJY" "RGN" ...
.. ..$ : chr [1:10] "PC1" "PC2" "PC3" "PC4" ...
- attr(*, "class")= chr "prcomp"
```

Application: NWSL Teams

```
# square root of eigenvalues  
pca$sdev
```

```
[1] 1.947118570 1.561340876 1.172896087 1.107625762  
[5] 0.849476477 0.502890350 0.365461414 0.205993340  
[9] 0.133840335 0.002521167
```

Application: NWSL Teams

```
# first three principal directions  
pca$rotation[, 1:3]
```

	PC1	PC2	PC3
possession_pct	0.3675695	-0.25738992	-0.07820678
goals	0.4218606	0.33869010	-0.09242574
assists	0.4152512	0.29222969	-0.09100741
pass_pct	0.3862252	-0.15782340	0.07035432
goal_conversion_pct	0.3368898	0.17920622	0.57472153
clean_sheets	0.2122449	-0.55909064	0.03992048
shot_accuracy	0.3389031	-0.07611405	0.04417828
shots_total	0.2224901	0.26001949	-0.66237229
goals_conceded	-0.1001183	0.51019377	0.36817595
tackled	-0.1799982	0.17416846	-0.25291961

How should we interpret the loadings for the first two principal directions?

Application: NWSL Teams

```
# original column standard deviations  
pca$scale
```

possession_pct	goals	assists
3.629634	5.710172	5.905827
pass_pct	goal_conversion_pct	clean_sheets
3.028963	3.185115	1.443376
shot_accuracy	shots_total	goals_conceded
3.481862	29.283877	5.944185
tackled		
32.577530		

Application: NWSL Teams

```
unscaled_pca <- prcomp(nwsl23_mat, scale = F)
unscaled_pca$rotation[, 1:3]
```

	PC1	PC2	PC3
possession_pct	0.001564728	-0.018339648	-0.12307260
goals	0.017632048	-0.145233100	-0.48496553
assists	0.001849595	-0.125199988	-0.54508628
pass_pct	-0.002864390	-0.015535596	-0.17098315
goal_conversion_pct	0.036297376	0.001347064	-0.38041324
clean_sheets	0.016597302	0.005625277	0.03499502
shot_accuracy	0.030876779	-0.043509933	-0.10711208
shots_total	-0.230507228	-0.954278241	0.12837166
goals_conceded	-0.064831781	0.040496550	-0.49884966
tackled	-0.969428672	0.220093406	-0.02377243

What changed? How can we interpret these loadings?

Essentially, if we do **not** scale the variables first, we are computing the eigenvectors of $X_c^\top X_c$, which is proportional to the sample covariance.

If we **do** scale the variables, we are computing the eigenvectors of the sample **correlation** matrix.

If our variables have differing sample variances, then the variables with larger variance will dominate the first principal components.

Application: NWSL Teams

```
biplot(pca)
```

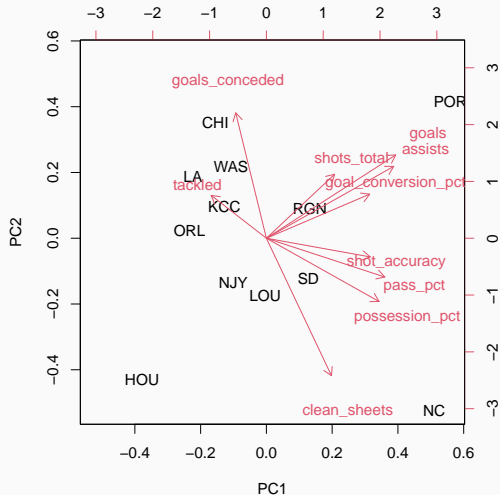


Figure 13: Biplot for NWSL 2023 team-level data

Biplots illustrate both:

- **principal components:** the positions of each observation in the rotated space (columns of U)
- **principal directions:** the columns of V contain **variable loadings** (the contribution of each variable to the PCs).

With which variables is the first principal component associated? What about the second principal component?

Application: NWSL Teams

```
plot(pca$rotation[,1], pca$rotation[,2], col = "white",  
     xlim = c(-.75, .75), ylim = c(-.75, .75),  
     xlab = "PC1", ylab = "PC2")  
text(pca$rotation[,1], pca$rotation[,2], rownames(pca$rotation))
```

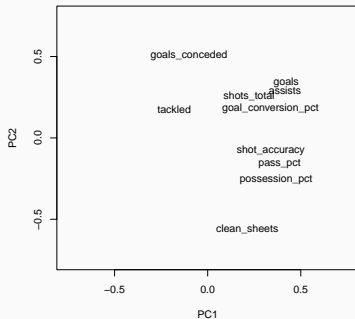


Figure 14: Variable loadings for first two PCs

Application: NWSL Teams

```
plot(pca$x[,1], pca$x[,2], col = "white",  
     xlim = c(-3.75, 3.75), ylim = c(-3.75, 3.75),  
     xlab = "PC1", ylab = "PC2")  
text(  
  pca$x[,1], pca$x[,2],  
  teams$team_abbreviation[match(nwsl23$team_id, teams$team_id)],  
)
```

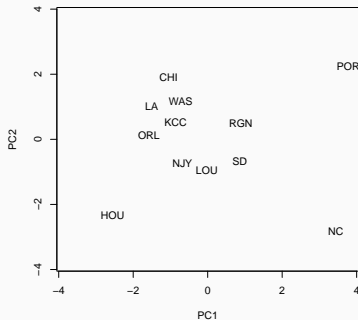


Figure 15: First two PCs for NWSL 2023 teams

Application: NWSL Teams

Pos	Team	[v•t•e]	Pld	W	D	L	GF	GA	GD	Pts	Qualification
1	San Diego Wave FC		22	11	4	7	31	22	+9	37	NWSL Shield, playoffs – semifinals
2	Portland Thorns FC		22	10	5	7	42	32	+10	35	Playoffs – semifinals
3	North Carolina Courage		22	9	6	7	29	22	+7	33	Playoffs – quarterfinals
4	OL Reign		22	9	5	8	29	24	+5	32	
5	Angel City FC		22	8	7	7	31	30	+1	31	
6	NJ/NY Gotham FC		22	8	7	7	25	24	+1	31	
7	Orlando Pride		22	10	1	11	27	28	−1	31	
8	Washington Spirit		22	7	9	6	26	29	−3	30	
9	Racing Louisville FC		22	6	9	7	25	24	+1	27	
10	Houston Dash		22	6	8	8	16	18	−2	26	
11	Kansas City Current		22	8	2	12	30	36	−6	26	
12	Chicago Red Stars		22	7	3	12	28	50	−22	24	

Figure 16: Standings for NWSL 2023 from Wikipedia

The first principal direction is associated with several variables about possession and scoring (goals and assists).

The second principal direction seems to have more to do with defense (goals conceded and clean sheets).

Application: NWSL Teams

```
var_explained = pca$sdev^2 / sum(pca$sdev^2)
ggplot(data = data.frame(x = c(1:10), y = var_explained),
       mapping = aes(x = x, y = y)) + geom_line() +
  xlab("PCs") + ylab("Variance Explained") + ggtitle("Scree Plot") +
  ylim(0, 1)
```

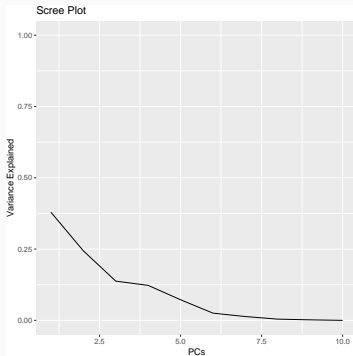


Figure 17: First two PCs for NWSL 2023 teams

Application: State-level characteristics

```
head(state.x77)
```

	Population	Income	Illiteracy	Life Exp	Murder
Alabama	3615	3624	2.1	69.05	15.1
Alaska	365	6315	1.5	69.31	11.3
Arizona	2212	4530	1.8	70.55	7.8
Arkansas	2110	3378	1.9	70.66	10.1
California	21198	5114	1.1	71.71	10.3
Colorado	2541	4884	0.7	72.06	6.8

	HS Grad	Frost	Area
Alabama	41.3	20	50708
Alaska	66.7	152	566432
Arizona	58.1	15	113417
Arkansas	39.9	65	51945
California	62.6	20	156361
Colorado	63.9	166	103766

Application: State-level characteristics

```
state_pca <- prcomp(state.x77, scale = T)
options(width = 60)
state_pca$rotation[, 1:3]
```

	PC1	PC2	PC3
Population	0.12642809	0.41087417	-0.65632546
Income	-0.29882991	0.51897884	-0.10035919
Illiteracy	0.46766917	0.05296872	0.07089849
Life Exp	-0.41161037	-0.08165611	-0.35993297
Murder	0.44425672	0.30694934	0.10846751
HS Grad	-0.42468442	0.29876662	0.04970850
Frost	-0.35741244	-0.15358409	0.38711447
Area	-0.03338461	0.58762446	0.51038499

Application: State-level characteristics

```
biplot(state_pca)
```

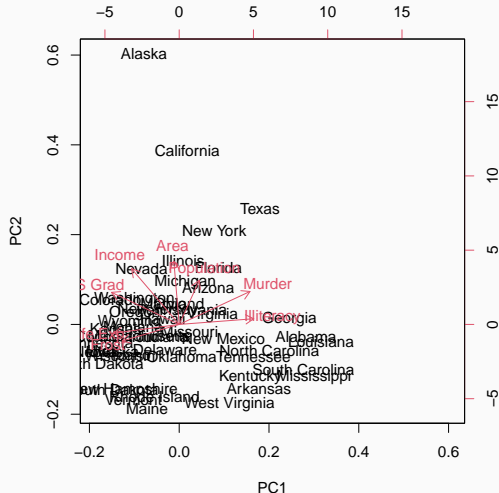


Figure 18: Biplot for state level characteristics, 1977

Application: State-level characteristics

```
plot(state_pca$rotation[,1], state_pca$rotation[,2], col = "white",  
     xlim = c(-.75, .75), ylim = c(-.75, .75),  
     xlab = "PC1", ylab = "PC2")  
text(state_pca$rotation[,1], state_pca$rotation[,2],  
     rownames(state_pca$rotation))
```

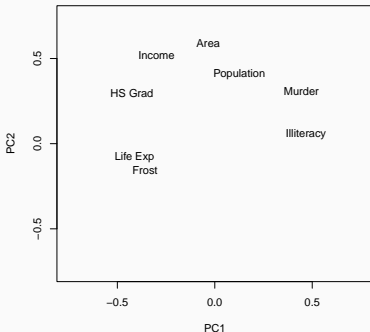


Figure 19: Variable loadings for PCA of state level characteristics, 1977

Application: State-level characteristics

```
plot(state_pca$x[,1], state_pca$x[,2], col = "white",
      xlim = c(-4.5, 4.5), ylim = c(-2, 6),
      xlab = "PC1", ylab = "PC2")
text(state_pca$x[,1], state_pca$x[,2], rownames(state.x77))
```

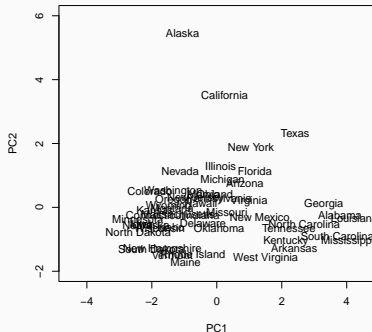


Figure 20: PCs for state level characteristics, 1977

Application: *The New York Times*

One way to turn documents into numerical data is to represent each document as a **bag of words**: a vector where each component represents the count of a particular word. These vectors are typically quite long and often **sparse**: many values are zero.

We can download a toy dataset from *the New York Times* [here](#) Shalizi (n.d.).

This dataset has 102 rows, and 4432 columns including the class label and the rest representing the counts for every distinct word that appears in at least one of the stories.

Application: *The New York Times*

```
head(nyt.frame)[, 1:6]
```

	class.labels		X.	X.d	X.nd	X.s
1	art	0.008706748	0.00000000	0	0.00000000	
2	art	0.005848328	0.00000000	0	0.00000000	
3	art	0.016035669	0.00000000	0	0.01140303	
4	art	0.026414939	0.00000000	0	0.00000000	
5	art	0.007285014	0.00000000	0	0.01100835	
6	art	0.002158439	0.03363435	0	0.03913930	

	X.th
1	0.009251444
2	0.000000000
3	0.000000000
4	0.000000000
5	0.000000000
6	0.000000000

Application: *The New York Times*

```
# drop class labels
nyt.pca <- prcomp(nyt.frame[, -1])
nyt.latent.sem <- nyt.pca$rotation
head(nyt.pca$rotation[, 1:3])
```

	PC1	PC2	PC3
X.	0.027008304	-0.005153922	-0.044675585
X.d	0.040733115	0.002834693	-0.026691733
X.nd	-0.006573117	0.004266679	-0.007905854
X.s	-0.022760270	0.036805485	0.038662704
X.th	0.035216761	0.014050485	-0.030373464
X.this	-0.004997333	0.013267267	0.005533896

Application: *The New York Times*

```
# Largest positive coordinates for the second PC  
signif(sort(nyt.latent.sem[, 2], decreasing = TRUE)[1:30], 2)
```

art	museum	images	artists	donations
0.150	0.120	0.095	0.092	0.075
museums	painting	tax	paintings	sculpture
0.073	0.073	0.070	0.065	0.060
gallery	sculptures	painted	white	patterns
0.055	0.051	0.050	0.050	0.047
artist	nature	service	decorative	feet
0.047	0.046	0.046	0.043	0.043
digital	statue	color	computer	paris
0.043	0.042	0.042	0.041	0.041
war	collections	diamond	stone	dealers
0.041	0.041	0.041	0.041	0.040

Application: *The New York Times*

```
# Largest negative coordinates for the second PC
```

```
signif(sort(nyt.latent.sem[, 2], decreasing = FALSE)[1:30], 2)
```

her	she	theater	opera
-0.220	-0.220	-0.160	-0.130
ms	i	hour	production
-0.130	-0.083	-0.081	-0.075
sang	festival	music	musical
-0.075	-0.074	-0.070	-0.070
songs	vocal	orchestra	la
-0.068	-0.067	-0.067	-0.065
singing	matinee	performance	band
-0.065	-0.061	-0.061	-0.060
awards	composers	says	my
-0.058	-0.058	-0.058	-0.056
im	play	broadway	singer
-0.056	-0.056	-0.055	-0.052
cooper	performances		
-0.051	-0.051		

Application: *The New York Times*

```
plot(  
  nyt.pca$x[, 1:2],  
  pch = ifelse(nyt.frame[, "class.labels"] == "music", "m", "a"),  
  col = ifelse(nyt.frame[, "class.labels"] == "music", "blue", "red")  
)
```

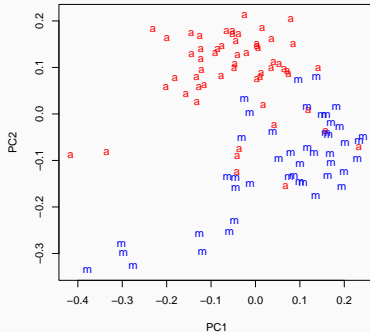
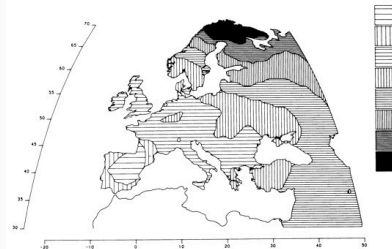
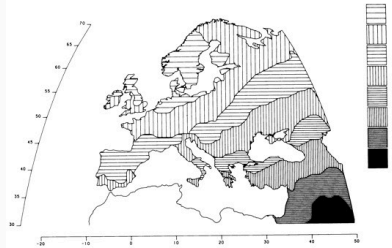


Figure 21: Projection of articles on the first two PCs.

Pitfalls of PCA

- It is common to try to interpret the principal components, but it's important to be cautious. We should be wary of "reifying" concepts.
- A key example comes from Cavalli-Sforza (1997), who describes a PCA with a data matrix where the rows represent locations and columns represent frequency of gene variants.

Cavalli-Sforza et al. (1997): Population migration from PCs?



Cavalli-Sforza et al. (1997): Population migration from PCs?

*"Hidden patterns in the geography of Europe shown by the first five principal components, explaining respectively 28%, 22%, 11%, 7%, and 5% of the total genetic variation for 95 classical polymorphisms. **The first component is almost superimposable to the archaeological dates of the spread of farming from the Middle East between 10,000 and 6,000 years ago.** The second principal component parallels a probable spread of Uralic people and/or languages to the northeast of Europe... "*

Shalizi (n.d.) reviews a paper by Novembre and Stephens (2008) that points out that these kinds of patterns are expected when carrying out PCA with **any** spatially correlated data.

Novembre and Stephens simulated data based on genetic diffusion processes, without any migration/population expansion and produced similar maps.

Overinterpretation

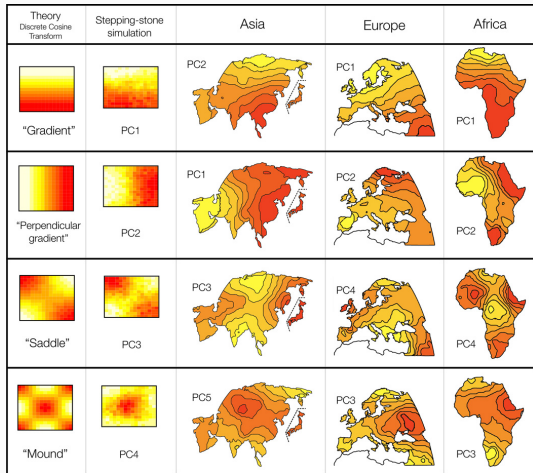


Figure 22: PCs based on simulated data with no migration

In other words, Novembre and Stephens do not disprove that migration happened, but they show that PCA of Cavalli-Sforza et al. doesn't provide strong evidence of the migration.

PCs must thus be interpreted with caution.

Linear projections

```
theta <- runif(100, min = 0, max = 2 * pi)
x <- cos(theta)
y <- sin(theta)
plot(x, y, pch = 16)
```

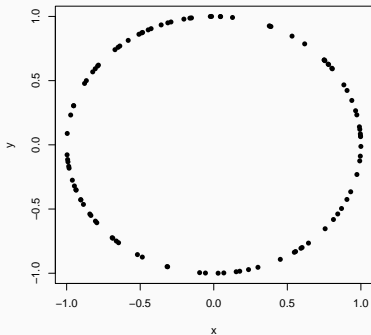


Figure 23: Circular data

Linear projections

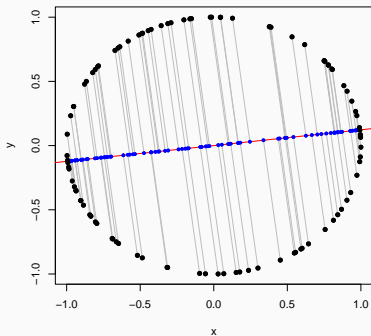


Figure 24: First PC and projection of circular data